

pg-flight-recorder

A Black Box for PostgreSQL

David A. Ventimiglia

Supabase

The Problem

"What was happening in my database?"

- 3 AM: Users complaining, dashboards red
- You wake up, check PostgreSQL... everything looks fine now
- **The incident ended. The evidence vanished.**

PostgreSQL forgets the past:

- `pg_stat_activity` shows **now, not then**
- No built-in history of wait events, locks, or queries

The Solution

Flight recorder concept: Always recording, ready when you need it


Like an airplane's black box - you don't think about it until you need it

- **Server-side** - runs inside PostgreSQL via pg_cron
- **Zero-config** - `psql -f install.sql` and done
- **Production-safe** - battle-tested at Supabase scale

What It Captures

Category	Metrics
Wait Events	What backends are waiting on
Active Queries	Running queries with session age
Locks	Lock contention and blocking chains
WAL & Checkpoints	Write-ahead log activity
I/O Timing	Storage latency by backend type
Top Queries	Slowest/most frequent queries
Table/Index Stats	Hotspots, bloat, unused indexes
Config Changes	Parameter drift detection

How It Works



Production Safe

Measured overhead: ~0.02% CPU

Metric	Value
Median collection	23ms
P95	31ms
Collections/day	480

Built-in protection:

- Load shedding at 70% connection utilization
 - Circuit breaker for slow collections
 - Auto-disable at 10GB storage
- 0% DDL blocking (tested: 202 operations)

Quick Demo: The Report

```
SELECT flight_recorder.report('1 hour');
```

Output includes:

- Anomalies detected (checkpoints, buffer pressure, locks)
- Wait event summary
- Lock contention events
- Long-running transactions
- Configuration changes
- Table hotspots & index efficiency

Paste into ChatGPT for instant analysis!

Real Investigation: Compare

```
SELECT * FROM flight_recorder.compare(  
    '2024-01-15 02:00', -- before incident  
    '2024-01-15 03:00'  -- during incident  
);
```

metric	before	during	delta
checkpoint_occurred	false	true	
bgw_buffers_backend_delta	0	12847	+12847
temp_bytes_delta	0	2.1 GB	+2.1GB

Diagnosis: Checkpoint during peak + backends writing directly to disk

Anomaly Detection

Auto-detects 12 issue types:

Anomaly	What It Means
CHECKPOINT_DURING_WINDOW	I/O spike from checkpoint
BUFFER_PRESSURE	Backends bypassing bgwriter
LOCK_CONTENTION	Sessions blocked on locks
XID_WRAPAROUND_RISK	Transaction ID approaching limit
IDLE_IN_TRANSACTION	Sessions blocking vacuum
DEAD_TUPLE_ACCUMULATION	Table bloat building up
REPLICATION_LAG_GROWING	Replica falling behind

Installation

```
psql -f install.sql
```

That's it. Collection starts automatically.

Requirements:

- PostgreSQL 15, 16, or 17
- pg_cron extension
- Optional: pg_stat_statements for query analysis

Uninstall:

```
psql -f uninstall.sql
```

Try It Today

github.com/supabase/pg-flight-recorder

```
git clone https://github.com/supabase/pg-flight-recorder
psql -f install.sql
# ... wait for an incident ...
SELECT flight_recorder.report('1 hour');
```

MIT Licensed | Open Source | Production-Proven

Questions?